

# Inference Latency vs. Number of Servers (Inception V3)

