

# Inference Latency vs. Number of Servers (ViT)

