

Inference Latency vs. Number of Servers (TinyBERT-6I)

