

# Inference Latency vs. Number of Servers (DistillBERT-base)

