

Inference Latency vs. Number of Servers (ViT-large)

