

# Inference Latency vs. Number of Servers (TinyBERT-6L)

